

Bruce Campell NCSU ST 534 Exam 2


16 November, 2017

1 sunspotz analysis

Consider the series sunspotz in the package astsa and answer the following questions:

(a) Plot sunspotz. Comment on the notable features of the time series?


```
rm(list = ls())  
library(astsa)  
data(sunspotz, package = "astsa")  
df <- sunspotz  
plot(ts(df))
```



From the documentation sunspotz. is the *biannual smoothed (12-month moving average)* number of sunspots from June 1749 to December 1978; $n = 459$. The format is **Time Series: Start = c(1749, 1) End = c(1978, 1) Frequency = 2** There appear to be two or more fundamental frequencies in the series. We note the appearance of short term and long term cyclical behavior. In regards to modelling with this data we would want to take into consideration that the period of observations is short relative to the evolution of the dynamical system generating the data, and that the technology for making observations may have evolved over the period of data collection.

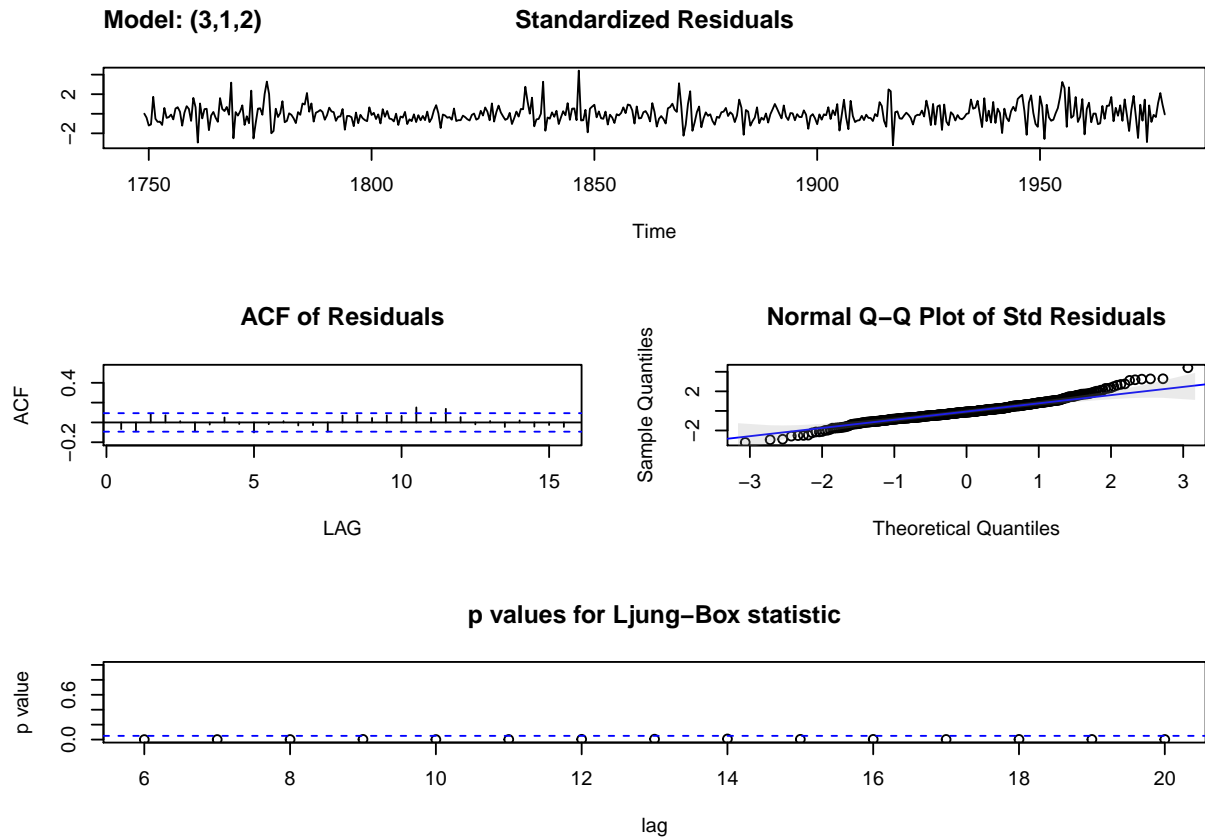
(b) Plot the ACF of sunspotz. Do you think it is appropriate to model this using a stationary ARMA process ?

```
invisible(acf2(df))
```



We note the presence of strong long range correlation. At first glance we would be hesitant to model this with a stationary ARMA process. We'd possibly perform a unit root test and consider modelling this data with an ARIMA process. For fun, we do some modelling in the temporal domain.

```
invisible(model <- sarima(df, 3, 1, 2))
```




```
model$tttable
```

We note some outlier in the the standardised residuals, heavy tails in the QQ plot, and that the individual Ljung-Box pp-vaules are all close to being significant. We also note a presense of a lag 11,12 acf component in the rsiduals. This is consistent with our intial observation that the data had short term and long term cyclic trends.

For more fun let's calculate the Ljung-Box-Pierce Q-statistic to check for systemic autocorrelation in the residuals.

```
n <- length(df)
H <- 20
r <- model$fit$residuals[1:H]
acf.residuals <- acf(r, H, main = "ACF of residuals")
```

ACF of residuals



```
sum.denominator <- n - seq(H, 1, by = -1)
r.s <- acf.residuals$acf^2/sum.denominator
Q <- n * (n + 2) * sum(r.s)
Q
```


```
## [1] 749.1462
```

We see based on the Q-statistic that we have significant correlation structure remaining in the residuals.

(c) Plot the periodogram of the sunspotz series using the spectrum command.

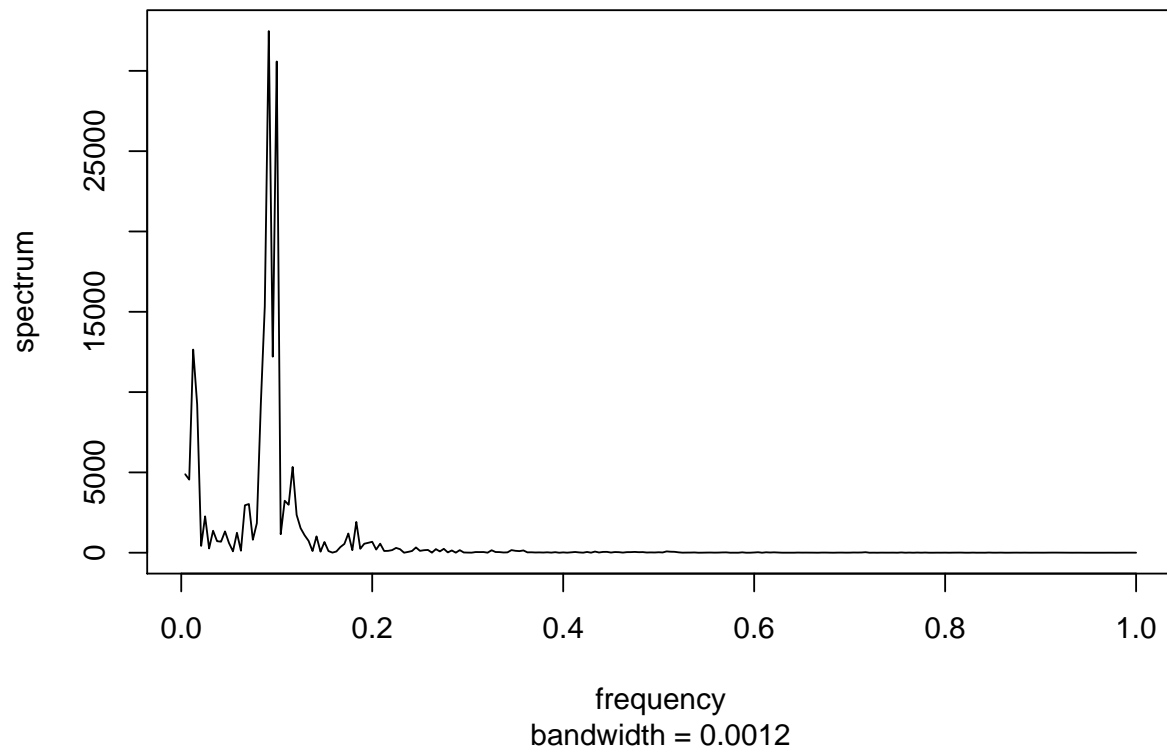
We used the definition of the periodogram and the fft to calculate $I(\omega_j) = |d(\omega_j)|^2$ taking into account that we had to normalize the fft result, take out the DC component, and only use half of the fft. We also plotted with the spec.pgram

```
df.periodogram <- Mod(fft(df - mean(df)))^2/length(df)
df.half <- ts(df.periodogram[1:length(df)/2])
plot(df.half, main = "Periodogram from fft")
```



```
df.periodogram.spec <- spec.pgram(df, taper = 0, log = "no", main = "Periodogram from sp
```

Periodogram from spec



What we have plotted is the raw periodogram. The raw periodogram is not a consistent estimator for the spectral density. We could optionally use a smoothing kernel in the `spec.pgram` function to obtain a consistent estimate.

(d) Find the maximum of the periodogram values and the frequency at which the

maximum occurs.

We're interested in the max and the other peaks to try and understand the frequency components in the signal. Here we extract the max. The `spec.pgram` method returns a spectrum object with the data in `freq`, `spec` lists.

```
max.spec.loc <- which.max(df.periodogram.spec$spec)
max.spec <- df.periodogram.spec$spec[max.spec.loc]
max.spec.freq <- df.periodogram.spec$freq[max.spec.loc]

pander(data.frame(max.spec.freq = max.spec.freq, max.spec = max.spec))
```


max.spec.freq	max.spec
0.09167	32478

We tried to fix up some code to find locations of the other peaks in the periodogram. The idea is to use overlapping intervals to find the local peaks via a spline maximization. It doesn't work yet :(

We did try to smooth the spectrum with a Daniell Kernel first. We'll keep that plot for reference. Note the 2 periodic components are highlighted better in the Daniell smoothed spectrum.

```
k = kernel("daniell", 4)
df.periodogram.spec.daniell <- spec.pgram(df, k, taper = 0, log = "no", main = "Periodogram from spec with Daniell kernel smoothing.")
```

Periodogram from spec with Daniell kernel smoothing.




```
# f <- function(x, q, d) spline(q, d, xout = x)$y x <-
# df.periodogram.spec.daniell$freq y <- df.periodogram.spec.daniell$spec nb
# <- 20 # choose number of intervals iv <- embed(seq(floor(min(x)),
# ceiling(max(x)), len = nb), 2)[,c(2,1)] # make overlapping intervals to
# avoid problems if the peak is close to # the ends of the intervals (two
# modes could be found in each interval) iv[-1,1] <- iv[nrow(iv),2] - 2 #
# The function 'f' is maximized at each of these intervals iv # the solution
# is a local maximum gr.thr <- 0.01 hes.thr <- 0.03i require('numDeriv')
# vals <- matrix(nrow = nrow(iv), ncol = 3) grd <- hes <- rep(NA,
# nrow(vals)) for (j in seq(1, nrow(iv))) { opt <- optimize(f = f, maximum =
# TRUE, interval = iv[j,], q = x, d = y) vals[j,1] <- opt$max vals[j,3] <-
```

```
# exp(opt$objj) grd[j] <- grad(func = f, x = vals[j,1], q = x, d = y) hes[j]
# <- hessian(func = f, x = vals[j,1], q = x, d = y) if (abs(grd[j]) < gr.thr
# && abs(hes[j]) > hes.thr) vals[j,2] <- 1 } # bin the peaks to avoid
# similar local maxima vals[,1] <- round(vals[,1], 2) if (anyNA(vals[,2])) {
# peaks <- unique(vals[-which(is.na(vals[,2])),1]) } else peaks <-
# unique(vals[,1]) plot(df.periodogram.spec$freq, df.periodogram.spec$spec,
# log = 'y', type = 'l') abline(v = peaks, lty = 2)
```

(e) Find an AR-based estimate of the spectral density of sunspotz using the command `spec.ar`.

First we note the caution from the documentaiton of `spec.ar` *Warning Some authors, for example Thomson (1990), warn strongly that AR spectra can be misleading.* We'll keep this in mind as we proceed.

```
df.spectrum.ar <- spec.ar(df, log = "no") #, log='no', ylim=c(0,.5), plot = TRUE)
```



This looks similar to the Danniell smmothed spectrum only smoother. The second peak seems less distinct in this spectra.

(f) Find the maximum of the estimated spectral density and the frequency at which

the maximum occurs.

```
max.spec.loc <- which.max(df.spectrum.ar$spec)
max.spec <- df.spectrum.ar$spec[max.spec.loc]
max.spec.freq <- df.spectrum.ar$freq[max.spec.loc]

pander(data.frame(max.spec.freq = max.spec.freq, max.spec = max.spec))
```

max.spec.freq	max.spec
0.09619	31906

(g) Find the two estimates of the period of the most prominent periodic component

of the sunspotz series using your answers from parts (d) and (f). What does it say about the Sunspot cycle?

I'm not sure what we're being asked in this part. The period is given by $\frac{1}{\omega}$ so the max of the ar estimated spectral density has a period of $\frac{1}{0.09619}$ 10.3960911

The max of the spectrum given by the fft is $\frac{1}{0.09167}$ 10.9086942 So the period is likely between 10-11 years, but it's not clear right now how we would statistically estimate this.

We can look at the frequency of the max of the Danniel smoothed spectrum for additional insight.

```
max.spec.loc <- which.max(df.periodogram.spec.danniel$spec)
max.spec <- df.periodogram.spec.danniel$spec[max.spec.loc]
max.spec.freq <- df.periodogram.spec.danniel$freq[max.spec.loc]
pander(data.frame(max.spec.freq = max.spec.freq, max.spec = max.spec))
```

max.spec.freq	max.spec
0.1	12511

We learned how to parametrically and non-parametrically estimate the spectrum $f(\omega_0)$ at a location ω_0 . Maybe we'll return to this and perform those estimates.

2

Suppose that $\{S_t\}$ and $\{N_t\}$ are zero mean SOS and independent time series with ACvFs $\gamma_1()$, γ_2 and spectral densities $f_1(\omega)$ and $f_2(\omega)$, respectively. Let

$$X_t = S_t + AS_{t-D} + N_t$$

where $A \in (0, \infty)$ is a constant and $D > 0$ is an integer.

(a) Find the ACvF of $\{X_t\}$.

$$f_x(\omega) = [1 + A^2 + 2A \cos(2\omega D)]f_1(\omega) + f_2(\omega)$$

Without loss of generality we can assume that $E[S_t] = 0$ and $E[N_t] = 0$ otherwise we'll subtract μ_S and μ_N from S_t and N_t and $\mu_S + A\mu_S + \mu_N$ from X_t . Now that we've mean corrected, we can write all of the ACvF functions in the reduced form.

$$\gamma_X = E[X_{t+h}X_t] = E[(S_{t+h} + AS_{t+h-D} + N_{t+h})(S_t + AS_{t-D} + N_t)]$$

Since the processes are independent the cross covariance terms drop out and we're left with

$$\gamma_X = E[X_{t+h}X_t] = E[S_{t+h}S_t] + A^2 E[S_{t-D+h}S_{t-D}] + E[N_{t+h}N_t] + A E[S_{t+h}S_{t-D}] + A E[S_{t-D+h}S_t]$$

Now we write this in terms of γ_S and γ_N

$$\gamma_X(h) = \gamma_S(h) + A^2 \gamma_S(h) + \gamma_N(h) + A \gamma_S(h+D) + A \gamma_S(h-D)$$

(b) Using part (a) or otherwise, show that the spectral density of $\{X_t\}$ is given by

Looking good so far, we can see where it's going. But we need to make an assumption regarding the summability of γ_S and γ_N . If we have absolute summability of the autocovariance functions we'll be able to express the spectral densities in terms of the values of $\gamma(h)$. We'll make this assumption and move on. The plan is that we'll express $f_X(\omega)$ in terms of the densities of S and N and invoke the uniqueness of the Fourier transform.

$$f_X(\omega) = \sum_{h=-\infty}^{h=\infty} \gamma_X(h) e^{-2\pi i \omega h} = \sum_{h=-\infty}^{h=\infty} [\gamma_S(h) + A^2 \gamma_S(h) + \gamma_N(h) + A \gamma_S(h+D) + A \gamma_S(h-D)] e^{-2\pi i \omega h}$$

Now we note that

$$A \sum_{h=-\infty}^{h=\infty} \gamma_S(h+D) e^{-2\pi i \omega h} = A e^{2\pi i \omega D} \sum_{h=-\infty}^{h=\infty} \gamma_S(h+D) e^{-2\pi i \omega (h+D)} = A e^{2\pi i \omega D} \sum_{(h+D)=-\infty}^{(h+D)=\infty} \gamma_S(h+D) e^{-2\pi i \omega (h+D)}$$

$$= A e^{2\pi i \omega D} \sum_{h=-\infty}^{h=\infty} \gamma_S(h) e^{-2\pi i \omega (h)}$$

We get a similar result for the term with $h - D$ except for a sign change. Substituting these expressions and those for $f_S(\omega)$ and $f_N(\omega)$ into our equation for $f_X(\omega)$ we have that

$$f_X(\omega) = f_S(\omega) + A^2 f_S(\omega) + f_N(\omega) + A e^{2\pi i \omega D} f_S(\omega) + A e^{-2\pi i \omega D} f_S(\omega)$$

Collecting terms and using a trig identity, we have that.


$$f_X(\omega) = [1 + A^2 + A \cos(2\pi i \omega D)] f_S(\omega) + f_N(\omega)$$

Uniqueness of the Fourier transform assures us that the expression on the RHS is indeed the spectral density of X .


3. Consider the unemployment time series `unemp` given in the package `astsa`.

(a) Plot the `unemp` series and its ACF upto lag 50.

```
rm(list = ls())
library(astsa)
data(unemp, package = "astsa")
df <- unemp
plot(ts(df))
```




```
invisible(acf2(df, 50))
```



(b) Fit a Seasonal ARIMA model of order $(012) \times (0, 1, 1)_{12}$ and present the R output.

```
invisible(model <- sarima(df, 0, 1, 2, 0, 1, 1, 12))
```




```
model$tttable
```

For lag 1 the Ljung-Box statistic shows significant correlation in the residuals.

Let's calculate the Ljung-Box-Pierce Q-statistic to check for systemic autocorrelation in the residuals. We'll extract the residuals and do the calculation by hand - there's an R function for this Box-Test that we've experimented with. we'll revisit this.

```
n <- length(df)
H <- 20
r <- model$fit$residuals[1:H]
acf.residuals <- acf(r, H, main = "ACF of residuals")
```

ACF of residuals



```
sum.denominator <- n - seq(H, 1, by = -1)
r.s <- acf.residuals$acf^2/sum.denominator
Q <- n * (n + 2) * sum(r.s)
Q
```

```
## [1] 506.2771
```

(c) Consider the Ljung-Box statistic p-value plot and the ACF of the residuals. What

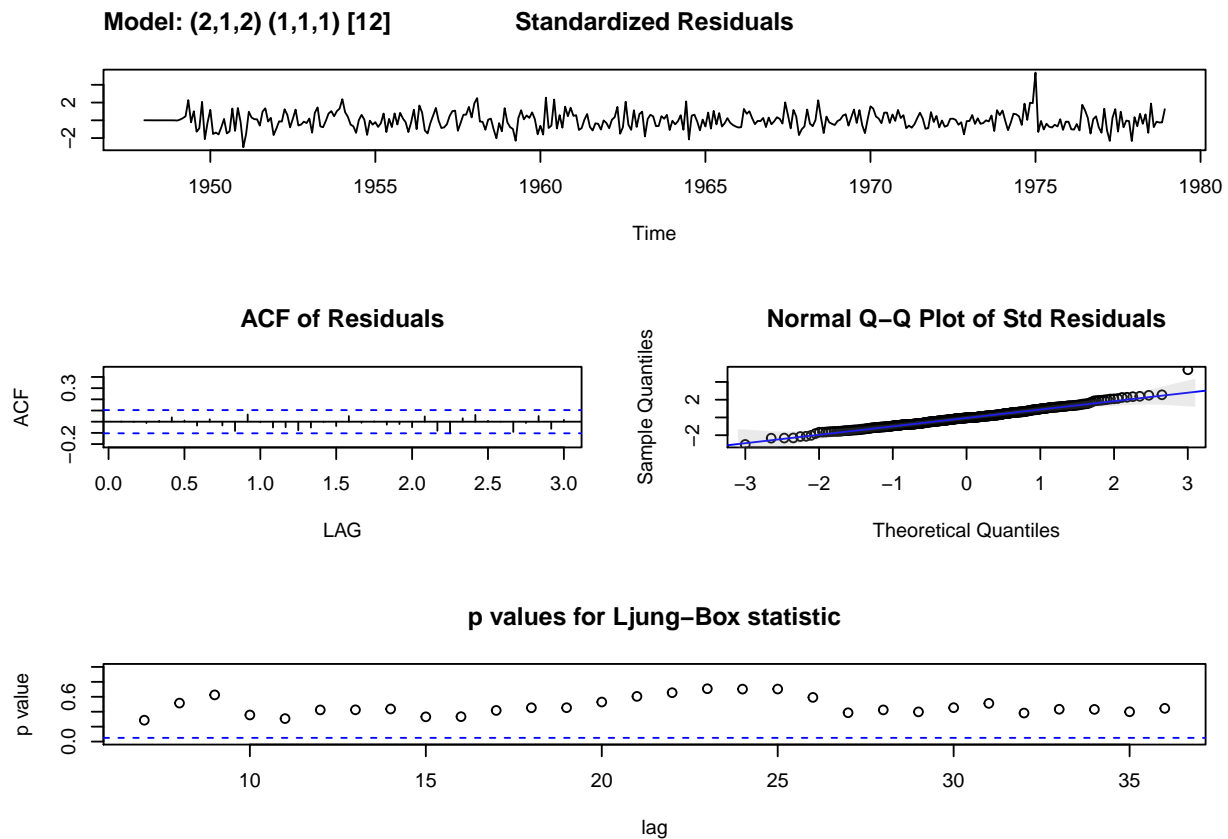
do these indicate about inadequacy of the fitted model ?

The Ljung-Box statistic shows nearly significant and autocorrelation in the residuals at some lags. We would look into increasing the AR order for our model to reduce this.

(d) Refine the SARIMA model based on your answer to the last part and present the

R output for the new SARIMA model. Indicate how the new model improves on the model in part (b).

```
invisible(model <- sarima(df, 2, 1, 2, 1, 1, 1, 12))
```




```
model$tttable
```

In looking at my calculation of the Ljung-Box-Pierce Q-statistic I'm suspicious of my code based on the plot above. REVISIT

```
n <- length(df)
H <- 20
r <- model$fit$residuals[1:H]
acf.residuals <- acf(r, H, main = "ACF of residuals")
```


ACF of residuals



```
sum.denominator <- n - seq(H, 1, by = -1)
r.s <- acf.residuals$acf^2/sum.denominator
Q <- n * (n + 2) * sum(r.s)
Q
```

```
## [1] 508.2607
```

(e) Use the fitted model in part (d) to forecast the next 12-months.